

**Full system  
simulation in  
gem5**



# What is Full System Simulation?

gem5's **Full System (FS) mode** simulates an entire computer system. This is in contrast to SE mode which uses the host OS.

# Basics of Booting Up a Real System in gem5

Unlike in SE mode where we can just provide a binary, in FS mode we need to provide much more information to boot up a real system. In particular we need:

1. A disk image containing the operating system and any necessary software or data. This disk image serves as the virtual hard drive for the simulated system.
2. A kernel binary compatible with the simulated architecture is needed to boot the operating system.

Beyond these essentials, you might need to provide other files like a bootloader, depending on the complexity of the simulation.

# How does gem5 know it's in FS mode or SE mode?

In `simulator.py` the root object is created, in which the `full_system` parameter is set to `True` or `False` depending on whether the full system is being simulated.

```
root = Root(
    full_system=(
        self._full_system
        if self._full_system is not None
        else self._board.is_fullsystem()
    ),
    board=self._board,
)
```

In `abstract_board.py`:

```
def is_fullsystem(self) -> bool:
    # ...
    if self._is_fs == None:
        raise Exception(
            "The workload for this board not ..."
        )
    return self._is_fs
```

# Ultimately determined by what `set_workload` function is called

In `kernel_disk_workload.py`

```
class KernelDiskWorkload:
    def set_kernel_disk_workload(
        self,
        kernel: KernelResource,
        disk_image: DiskImageResource,
        bootloader: Optional[BootloaderResource] = None,
        disk_device: Optional[str] = None,
        ...
    ) -> None:
        ...
        self._set_fullsystem(True)
```

## These `set_workload` classes mixin with the boards

```
class X86Board(AbstractSystemBoard, KernelDiskWorkload):  
    ...
```

So, in short, **the `set_workload` function called determines whether the simulation is in FS mode or SE mode.**

# Accessing the running system

In FS mode, you can access the running system using the **serial console**.

You can use the `m5term` utility to connect to the serial console.

Build `m5term` using the following command:

```
cd gem5/util/term  
make
```

Now you have a program called `m5term` in the `gem5/util/term` directory.

When you run `gem5`, it will print `Listening for connections on port 3456` when it is ready to accept connections from `m5term`.

You can then run:

```
./m5term localhost 3456
```

## Exercise: explore a running simulation

Create a script that boot Ubuntu 22.04 and run some commands.

Use the x86 ISA and the `CPUTypes.KVM` CPU type.  
We'll use KVM so that it is fast.

Use the "x86-ubuntu-22.04-boot-with-systemd" workload from gem5-resources. (This has already been downloaded for you.)

You can run in interactive mode by adding the following line after `board.set_workload`.

```
board.append_kernel_arg("interactive=true")
```

**IMPORTANT:** See the next slide for two important notes before you run.



# Important notes

## Note 1

Make sure that you have the following line in your script. This is required for running in codespaces, but may not be required on your local machine.

```
for proc in processor.cores:  
    proc.core.usePerf = False
```

## Note 2

The disk image that you're using has some exit events that will execute by default. You need to make sure to handle them. Here is some example code.

```
def exit_event_handler():  
    print("first exit event: Kernel booted")  
    yield False  
    print("second exit event: In after boot")  
    yield False  
    print("third exit event: After run script")  
    yield True  
simulator = Simulator(  
    board=board,  
    on_exit_event={  
        ExitEvent.EXIT: exit_event_handler(),  
    },  
)  
simulator.run()
```

# Using "readfile"

gem5 has a utility called `readfile` that can be used to read files from the simulated system. This can be used to execute a command automatically after boot.

You can see how it's used in the default disks by looking at the `after_boot.sh` file in the disk image.

See `after_boot.sh` in the gem5-resources.

(Or at [https://github.com/gem5/gem5-resources/blob/stable/src/ubuntu-generic-diskimages/files/x86/after\\_boot.sh](https://github.com/gem5/gem5-resources/blob/stable/src/ubuntu-generic-diskimages/files/x86/after_boot.sh))

# Overriding kernel disk workload

You can override the inputs to the workload by calling `set_kernel_disk_workload` directly.

Look at the [details of the workload you have been using](#) on gem5-resources to get started.

The key parts are the `resources` and `additional_params` sections.

Use these parameters to fill in the arguments on `set_kernel_disk_workload`.

# Running your own command

Now, you should have something like below, and you can override the `readfile_contents` argument.

```
board.set_kernel_disk_workload(  
    kernel = obtain_resource("x86-linux-kernel-5.4.0-105-generic"),  
    disk_image = obtain_resource("x86-ubuntu-22.04-img"),  
    kernel_args = ["earlyprintk=ttyS0", "console=ttyS0", "lpj=7999923", "root=/dev/sda2"],  
    readfile_contents = "echo 'Hello, world!'; sleep 5",  
)
```

We add `sleep 5` to keep the simulation running so the serial console will be flushed.

Now, when you run the simulation, you should see the output of the `echo` command in the terminal output. (You will want to remove the `board.append_kernel_arg("interactive=true")` line.)

You can also specify a file with the commands via `readfile`.

# Creating your own disk images

# Creating disk images using Packer and QEMU

To create a generic Ubuntu disk image that we can use in gem5, we will use:

- Packer: This will automate the disk image creation process.
- QEMU: We will use a QEMU plugin in Packer to actually create the disk image.
- Ubuntu autoinstall: We will use autoinstall to automate the Ubuntu install process.

gem5 resources already has code that can create a generic Ubuntu image using the aforementioned method.

- Path to code: `gem5-resources/src/x86-ubuntu`

Let's go through the important parts of the creation process.



# Getting the ISO and the user-data file

As we are using Ubuntu autoinstall, we need a live server install ISO.

- This can be found online from the Ubuntu website: [iso](#)

We also need the user-data file that will tell Ubuntu autoinstall how to install Ubuntu.

- The user-data file on gem5-resources specifies all default options with a minimal server installation.

# How to get our own user-data file

To get a user-data file from scratch, you need to install Ubuntu on a machine.

- Post-installation, we can retrieve the `autoinstall-user-data` from `/var/log/installer/autoinstall-user-data` after the system's first reboot.

You can install Ubuntu on your own VM and get the user-data file.



# Using QEMU to get the user-data file

We can also use QEMU to install Ubuntu and get the aforementioned file.

- First, we need to create an empty disk image in QEMU with the command: `qemu-img create -f raw ubuntu-22.04.2.raw 5G`
- Then we use QEMU to boot the diskimage:

```
qemu-system-x86_64 -m 2G \  
    -cdrom ubuntu-22.04.2-live-server-amd64.iso \  
    -boot d -drive file=ubuntu-22.04.2.raw,format=raw \  
    -enable-kvm -cpu host -smp 2 -net nic \  
    -net user,hostfwd=tcp::2222-:22
```

After installing Ubuntu, we can use ssh to get the user-data file.

# Important parts of the Packer script

Let's go over the Packer file.

- **bootcommand:**

```
"e<wait>",  
"<down><down><down>",  
"<end><bs><bs><bs><bs><wait>",  
"autoinstall ds=nocloud-net\\;s=http://{{ .HTTPIP }}:{{ .HTTPPort }}/ ---<wait>",  
"<f10><wait>"
```

This boot command opens the GRUB menu to edit the boot command, then removes the `---` and adds the autoinstall command.

- **http\_directory:** This directory points to the directory with the user-data file and an empty file named meta-data. These files are used to install Ubuntu.

## Important parts of the Packer script (Conti.)

- **qemu\_args:** We need to provide Packer with the QEMU arguments we will be using to boot the image.
  - For example, the QEMU command that the Packer script will use will be:

```
qemu-system-x86_64 -vnc 127.0.0.1:32 -m 8192M \  
-device virtio-net,netdev=user.0 -cpu host \  
-display none -boot c -smp 4 \  
-drive file=<Path/to/image>,cache=writeback,discard=ignore,format=raw \  
-machine type=pc,accel=kvm -netdev user,id=user.0,hostfwd=tcp::3873-:22
```

- **File provisioners:** These commands allow us to move files from the host machine to the QEMU image.
- **Shell provisioner:** This allows us to run bash scripts that can run the post installation commands.

# Extending the base ubuntu image

Update the [x86-ubuntu.pkr.hcl](#) file.

The general structure of the Packer file would be the same but with a few key changes:

- We will now add an argument in the `source "qemu" "initialize"` block.
  - `diskimage = true` : This will let Packer know that we are using a base disk image and not an iso from which we will install Ubuntu.
- Remove the `http_directory = "http"` directory as we no longer need to use autoinstall.
- Change the `iso_checksum` and `iso_urls` to that of our base image.

# Step 1: Downloading the base Ubuntu image

Let's get the base Ubuntu 24.04 image from gem5 resources and unzip it.

```
wget https://storage.googleapis.com/dist.gem5.org/dist/develop/images/x86/ubuntu-24-04/x86-ubuntu-24-04.gz  
gzip -d x86-ubuntu-24-04.gz
```

`iso_checksum` is the `sha256sum` of the iso file that we are using. To get the `sha256sum` run the following in the linux terminal.

```
sha256sum ./x86-ubuntu-24-04.gz
```

## Step 2: Updating the packerfile

- **Update the file and shell provisioners:** Let's remove the file provisioners as we don't need to transfer the files again.
- **Boot command:** As we are not installing Ubuntu, we can write the commands to login along with any other commands we need (e.g. setting up network or ssh). Let's update the boot command to login and enable network:

```
"<wait30>",  
"gem5<enter><wait>",  
"12345<enter><wait>",  
"sudo mv /etc/netplan/50-cloud-init.yaml.bak /etc/netplan/50-cloud-init.yaml<enter><wait>",  
"12345<enter><wait>",  
"sudo netplan apply<enter><wait>",  
"<wait>"
```

## Step 3: Changes to the post installation script

For this post installation script we need to get the dependencies and build the GAPBS benchmarks.

Add this to the [post-installation.sh](#) script

```
git clone https://github.com/sbeamer/gapbs
cd gapbs
make
```

Let's run the Packer script and use this disk image in gem5!

```
cd /workspaces/2024/materials/04-Advanced-using-gem5/07-full-system
x86-ubuntu-gapbs/build.sh
```

## Step 4: Let's use our built disk image in gem5

Let's add the md5sum and the path to our [local JSON](#).

Let's run the [gem5 GAPBS config](#).

```
GEM5_RESOURCE_JSON_APPEND=./completed/local-gapbs-resource.json gem5 x86-fs-gapbs-kvm-run.py
```

This script should run the bfs benchmark.

